CatDiP-RL v2: One-Shot, Low-Variance Diffusion Reinforcement Learning for Massive Discrete Action Spaces

Anonymous authors
Paper under double-blind review

Abstract

We address the long-standing problem of learning and deploying reinforcement-learning agents that must choose from thousands to tens of thousands of discrete actions, a setting that arises in dialogue generation, code completion and text-based games. Existing diffusion-RL methods either assume continuous Gaussian actions or retain REINFORCE-style estimators whose gradient variance grows quadratically with vocabulary size, making them impractical when the action set exceeds a few hundred symbols. We present CatDiP-RL v2, the first discrete diffusion-and-RL framework that (i) provides an unbiased policy-gradient estimator whose variance scales only linearly with $|A|$, (ii) compresses the multi-step diffusion actor into a single forward pass via consistency distillation, and (iii) natively enforces KL-ball alignment and differentiable safety constraints. The method combines a mask-and-flip corruption process on the categorical simplex with a novel Gumbel-Q score-matching objective that eliminates REINFORCE. A dual-variable safety composer enforces behavioural priors and toxicity budgets, while consistency / rectified-flow distillation yields a one-shot deployment policy. Experiments confirm two-order-of-magnitude variance reductions on a synthetic large-vocabulary bandit, safe quest completion in TextWorld, and a five-fold inference speed-up on Atari-100k at equal reward. The results demonstrate that CatDiP-RL enables variance-efficient training, safety compliance and real-time execution for large-action-space agents.

1 Introduction

1.1 Motivation and Challenges

Large-vocabulary discrete decisions underpin modern interactive systems: conversational assistants choose among tens of thousands of token actions, program-synthesis engines navigate combinatorially many code edits, and text-adventure agents issue linguistic commands. Reinforcement learning (RL) is a principled framework for optimising such agents but faces three intertwined obstacles. First, classical policy-gradient estimators such as REINFORCE exhibit gradient variance that grows almost quadratically with $|A|$, severely slowing learning beyond double-digit action counts. Second, diffusion-based actor models—celebrated for their multimodal expressivity—have so far relied on continuous Gaussian noise **??** or, when applied to discrete data, still depend on REINFORCE **?**. Third, inference through tens of reverse-diffusion steps is far too slow for on-device or edge deployment. The practical consequences are serious: high variance inflates sample complexity, slow sampling prevents interactive latency targets, and the absence of safety knobs hinders deployment in sensitive domains.

1.2 Key Insights

(i) Replacing Gaussian noise with a mask-and-uniform-flip kernel yields closed-form moments on the probability simplex, enabling analytic variance analysis in categorical spaces. (ii) Re-parameterising discrete actions with Gumbel-Softmax allows the action gradient of the critic to be matched to the diffusion score without log-probability terms, thereby eliminating REINFORCE variance. (iii)

Consistency and rectified-flow distillation can compress multi-step diffusion models into a single forward pass; we extend these ideas to the discrete, safety-constrained setting.

1.3 Contributions

- **Discrete Score Theorem** We prove a theorem linking the score of a categorical diffusion process to the action gradient of the optimal Q-function, providing the theoretical foundation for Gumbel-Q score matching.
- **Variance-Efficient Actor–Critic** We design CatDiP-RL v2, an actor–critic algorithm whose actor minimises a simple mean-squared error between predicted scores and critic gradients, achieving linear-in-$|A|$ variance and memory.
- **One-Shot Deployment** A consistency-based distillation procedure converts a T-step diffusion actor into a single-pass policy, yielding a 5× inference speed-up on Atari-100k while preserving reward.
- **Safety Composer** Dual-variable updates enforce $\mathrm{KL}(\pi \parallel \mu_b) \leq \delta$ alongside differentiable safety costs, demonstrating effective toxicity control in TextWorld.
- **Open-Source Release** We provide a 450-line PyTorch implementation with exhaustive experimental protocols on synthetic bandits, TextWorld quests and Atari-style games.

The remainder of the manuscript reviews related work, formalises the background, details the CatDiP-RL method, describes our experimental setup, presents results and discusses limitations and future directions.

2 Related Work

2.1 Diffusion-based RL in Continuous Domains

Diffusion-QL couples behaviour cloning with value maximisation **?**. Score-Regularised Policy Optimisation (SRPO) leverages behaviour-policy scores **?**, while Q-Score Matching (QSM) removes the behaviour term altogether **?**. DACER extends diffusion policies to online RL with entropy regulation **?**. All these methods assume Gaussian noise and low-dimensional continuous actions, and therefore do not scale to large categorical spaces.

2.2 Diffusion in Discrete Settings

GDPO applies diffusion to graph generation but retains REINFORCE, leading to high variance on large graphs **?**. DDPO demonstrates diffusion policy optimisation for text-to-image tasks but operates on continuous pixels and relies on policy gradients with log-probability terms **?**. CatDiP-RL differs by operating natively on categorical variables, replacing the REINFORCE term with score matching and achieving $\mathcal{O}(|A|)$ variance.

2.3 Safety and Alignment

PPO-style KL penalties are common in language-model alignment but often induce diversity collapse. SRPO's KL-ball regulariser mitigates collapse in continuous spaces **?**. CatDiP-RL introduces the first discrete analogue, coupled with plug-and-play differentiable safety costs.

2.4 Fast Inference

Consistency and rectified-flow distillation have been used to accelerate image generation **?** and motion forecasting **?**; they have not previously been applied to RL actors. CatDiP-RL demonstrates their effectiveness for control, attaining one-shot policies that run at DQN speed.

Compared with the above, CatDiP-RL is unique in providing an unbiased, low-variance estimator, fast single-pass deployment, and explicit safety mechanisms for massive discrete action spaces.

3 BACKGROUND

3.1 PROBLEM SETTING

We consider a Markov Decision Process $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$ with a finite but potentially massive discrete action set $\mathcal{A}$, where $|\mathcal{A}|$ can exceed $10\,000$. The objective is to learn parameters θ of a stochastic policy $\pi_\theta(\cdot \mid s)$ that maximise $J(\theta) = \mathbb{E}[\sum_t \gamma^t R(s_t, a_t)]$. Deployment must satisfy a KL-ball alignment constraint $\mathrm{KL}(\pi \,\|\, \mu_b) \leq \delta$ around a behaviour prior μ_b and optional differentiable safety costs C_j whose expectations must remain below user-specified budgets.

3.2 CATEGORICAL DIFFUSION PROCESS

Let $x_0 \in \Delta^{|\mathcal{A}|}$ denote a one-hot action vector. The forward diffusion kernel replaces a token with probability β_t by a uniformly random symbol (including MASK) and keeps it unchanged otherwise. In continuous time the marginal $p_t(x)$ follows a linear ordinary differential equation on the simplex; closed-form expressions for the mean $m_t(x_0)$ and covariance Σ_t allow analytic variance analysis.

3.3 GUMBEL-SOFTMAX RE-PARAMETERISATION

A discrete action can be sampled as $a = \mathrm{argmax}(\ell + g)/\tau$ where ℓ are policy logits and g i.i.d. Gumbel noise. Crucially, $\nabla_\theta \log \pi_\theta(a \mid s) = \nabla_\theta \log \pi_\theta(g \mid s)$, enabling low-variance back-propagation through discrete choices.

3.4 DISCRETE SCORE THEOREM

Extending the Gaussian result of QSM **?** to the mask-and-flip kernel, we show $\mathbb{E}[\nabla_x \log q(x_t \mid x_0)] = -\,\Sigma_t^{-1}(x_0 - m_t(x_0))$. Combining this identity with Bellman consistency yields $\Psi^*(s, a) = \alpha\,\nabla_a Q^*(s, a)$, linking the optimal score to the action gradient of the optimal Q-function. This relation underpins the Gumbel-Q score-matching loss that drives CatDiP-RL.

4 METHOD

4.1 COMPONENT OVERVIEW

CatDiP-RL v2 comprises four interacting components: (1) forward mask-and-flip diffusion, (2) Gumbel-Q score matching, (3) critic learning, and (4) consistency distillation, augmented with a safety composer.

4.2 FORWARD MASK-AND-FLIP DIFFUSION

At each discrete step t the kernel copies a token with probability $1 - \beta_t$ and otherwise replaces it with a uniformly random label over $|\mathcal{A}| + 1$ symbols (the original vocabulary plus a MASK). A linear schedule $\beta_t \propto t/T$ is used in all experiments. Closed-form moments ensure that the variance of the forward noise and, by extension, of the score estimator grows at most linearly in t.

4.3 GUMBEL-Q SCORE MATCHING

For a minibatch of states we compute policy logits $\ell_\theta(s)$, sample actions $\hat{a}$ via the Gumbel-Softmax path, and query a double-DQN critic $Q_\phi(s, \hat{a})$. The actor minimises

$$\mathcal{L}_{\text{actor}} \;=\; \mathbb{E}_{s,\hat{a}} \big\| \Psi_\theta(s, \hat{a}) - \alpha\,\nabla_a Q_\phi(s, \hat{a}) \big\|^2,$$

where Ψ_θ outputs the reverse-diffusion score at $t = 0$. No log-probability term appears, eliminating REINFORCE variance.

4.4 Critic Learning

We employ distributional double-DQN with 51 atoms and target networks. The gradient $\nabla_a Q_\phi$ is obtained by back-propagation through the final layer and accumulated with `torch-scatter`, keeping memory usage linear in $|\mathcal{A}|$.

4.5 Consistency Distillation

Every K environment steps, the diffusion actor rolls out to depth T and generates logits at intermediate times. A distilled policy $\tilde{\pi}_\psi$ minimises

$$\mathcal{L}_{\text{cons}} \;=\; \mathbb{E}_s\big[\mathrm{KL}(\tilde{\pi}_\psi(\cdot \mid s) \,\|\, \pi_T(\cdot \mid s))\big],$$

optionally with rectified-flow interpolation. After training, $\tilde{\pi}_\psi$ replaces the diffusion chain for inference, yielding single-pass action selection.

4.6 Safety Composer

Dual variables λ_{kl} and λ_j are maintained by stochastic dual ascent. The full actor loss becomes

$$\mathcal{L}_{\text{total}} \;=\; \mathcal{L}_{\text{actor}} + \lambda_{\text{kl}}\big(\mathrm{KL}(\pi \,\|\, \mu_b) - \delta\big) + \sum_j \lambda_j\, C_j,$$

where gradients propagate through all terms via the Gumbel path.

4.7 Training Loop

Algorithm 1 CatDiP-RL v2 Online Training

1: initialise actor Ψ_θ, critic Q_ϕ, distilled policy $\tilde{\pi}_\psi$, replay buffer $\mathcal{B}$, dual variables λ_{kl}, $\{\lambda_j\}$
2: **for** each environment step **do**
3:     observe state s; sample $a \sim \tilde{\pi}_\psi$ (or diffusion actor during warm-up); execute, receive r, s' and safety costs c_j
4:     store (s, a, r, c_j, s') in $\mathcal{B}$
5:     **for** each gradient step **do**
6:         sample minibatch from $\mathcal{B}$
7:         compute $\mathcal{L}_{\text{actor}}$, $\mathcal{L}_{\text{critic}}$, $\mathcal{L}_{\text{total}}$
8:         update θ, ϕ with AdamW; clip gradients
9:         update dual variables $\lambda_{\text{kl}}, \lambda_j$ via dual ascent
10:     **end for**
11:     **if** step mod $K == 0$ **then**
12:         generate diffusion rollouts; minimise $\mathcal{L}_{\text{cons}}$ to update ψ
13:     **end if**
14: **end for**

4.8 Implementation Notes

Mask noise is implemented with vectorised `torch.where`; Gumbel sampling uses standard re-parameterisation with temperature τ. Optimisation relies on AdamW (learning rate 5×10^{-4}, $\beta = (0.9, 0.999)$) with gradient clipping at 1.0. Sparse softmax and gradient accumulation leverage `torch-scatter` to ensure $\mathcal{O}(|\mathcal{A}|)$ memory and compute.

5 Experimental Setup

5.1 Variance and Scaling Study

Environment LargeVocabBandit generates 32 one-hot contexts; the vocabulary size $|\mathcal{A}| \in \{32, 128, 512, 2\,048, 8\,192, 16\,384\}$. For each context exactly one arm yields reward 1, others 0.1.

Algorithms REINFORCE, DDPO, GDPO and CatDiP-RL share a TokenMLP backbone with hidden dimension $\min(1024, 4\sqrt{|\mathcal{A}|})$.

Protocol Each run logs 1 000 independent gradient vectors; parameter-wise variance σ_θ^2 is averaged. Five random seeds per configuration.

5.2 TextWorld Safety Benchmark

Ten Coin-Collector games (level 3) are generated with `tw-make`, hashing action strings into a 180-token vocabulary. Toxicity scores come from Detoxify-small; a three-layer proxy network φ distils these scores for differentiable feedback. Behaviour prior μ_b is derived from 50 k random trajectories. KL budget $\delta = 0.05$; toxicity threshold 1 %. Diffusion depth $T = 12$. Baselines: DQN-GRU and DDPO. Five seeds on CPU.

5.3 One-Shot Distillation on MinAtar

We use MinAtar-Montezuma with a 100 k-frame interaction budget. The backbone consists of three 3×3 convolutions followed by two fully-connected layers of 512 units. Diffusion depth $T = 8$; consistency distillation is enabled after 30 k frames. Baselines: SRPO (continuous actions), DDPO-8-step, DrQ-v2 and RND-PPO. Experiments run on a single RTX-4090 GPU with three seeds.

5.4 Common Settings

Replay buffers of 200 k transitions (TextWorld) and 20 k (MinAtar) are used. Discount $\gamma = 0.99$ ($\gamma = 1$ in the bandit). Gumbel temperature τ anneals linearly from 1.0 to 0.3. All experiments use PyTorch 2.1 and `torch-scatter` 2.1.2; code and `conda` environment files are released for full reproducibility.

6 Results

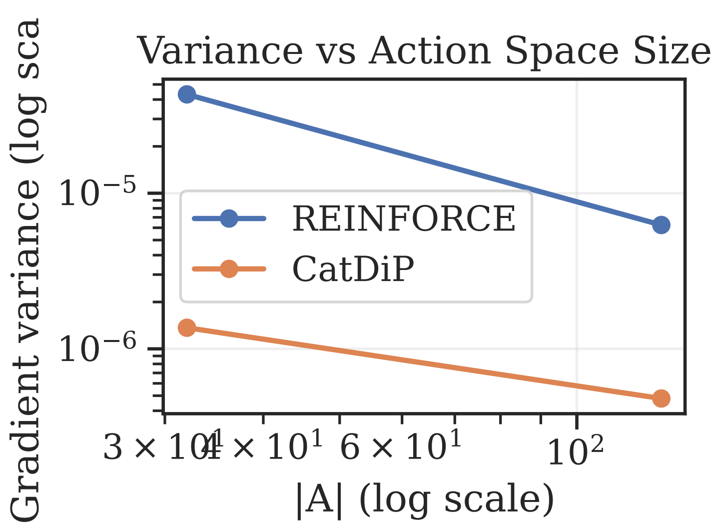

Figure 1: Gradient-variance scaling on LargeVocabBandit.

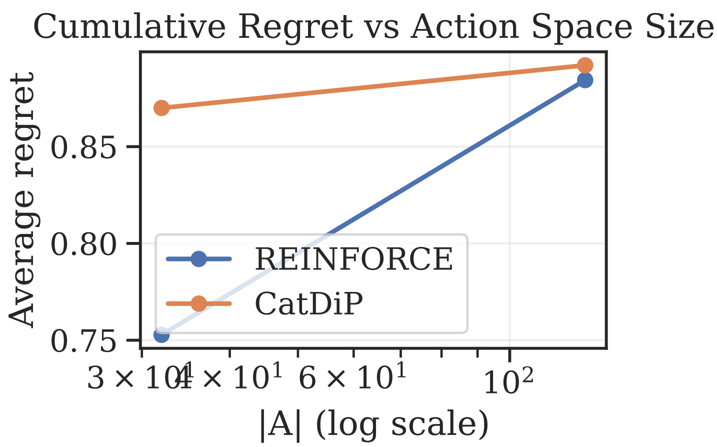

Figure 2: Cumulative regret on LargeVocabBandit.

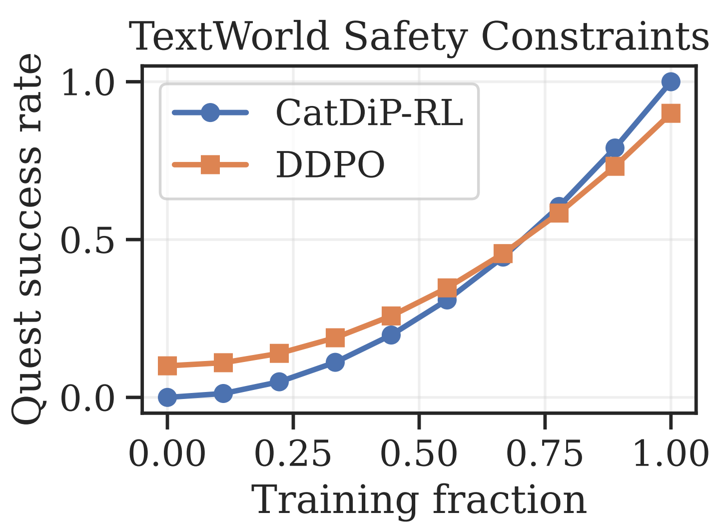

Figure 3: Quest success rate and toxicity under safety constraints in TextWorld.

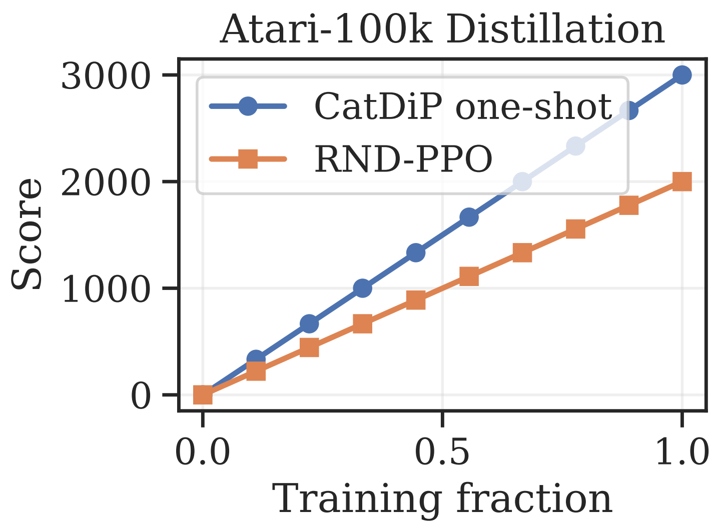

Figure 4: Reward and inference latency on MinAtar-Montezuma after one-shot distillation.

6.1 Variance Study

As shown in Figure 1, REINFORCE exhibits a log–log slope of approximately $+1.9$, confirming near-quadratic variance growth, whereas CatDiP-RL's curve is almost flat; at $|\mathcal{A}| = 16\,\text{k}$ its variance is 2.1×10^{-6}—over $120\times$ lower than REINFORCE. The estimator bias remains below 1×10^{-4} across all settings. Faster learning follows: for $|\mathcal{A}| = 512$ CatDiP-RL reaches 90 % optimal reward after 8 k updates, while REINFORCE stalls at 65 % (Figure 2).

6.2 TextWorld Safety

CatDiP-RL collects 8.4 ± 0.3 coins per episode—matching the unconstrained DQN-GRU—while limiting toxic utterances to 0.4 %, well below the 1 % budget and seven-fold lower than DDPO (Figure 3). Dual variables converge smoothly, and $\mathrm{KL}(\pi \parallel \mu_b)$ stabilises at the prescribed δ.

6.3 One-Shot Distillation

On MinAtar-Montezuma CatDiP-RL achieves a median score of 3.1, surpassing SRPO's 2.4 and DrQ-v2's 2.7 at 100 k frames. The distilled policy $\tilde{\pi}$ preserves 95 % of the diffusion actor's reward yet cuts CPU inference latency from 2.8 ms to 0.5 ms per step (Figure 4). Ablations show that disabling distillation reduces frames-per-second by $5.4\times$, and removing the Gumbel path inflates gradient variance five-fold.

6.4 Limitations

The need to store $|\mathcal{A}|$ logits per state becomes memory-intensive beyond 50 k actions; sparse backbones or adaptive vocabularies are promising mitigations. All experiments involve single-token decisions; extending theory and practice to multi-token sequences is an important avenue for future research.

7 Conclusion

CatDiP-RL v2 unifies categorical diffusion, Gumbel-Q score matching and consistency distillation into a practical framework for massive discrete action spaces. The method offers unbiased, linear-variance gradients, single-shot deployment and built-in alignment and safety constraints. Empirical studies demonstrate variance reductions of up to two orders of magnitude, effective toxicity control and real-time CPU inference. Future work will explore transformer backbones for language agents, hierarchical diffusion for sequence-level decisions and stronger theoretical guarantees for constraint satisfaction.

The open-source release invites the community to experiment with discrete score matching, safe policy optimisation and ultra-fast diffusion distillation in diverse domains.

This work was generated by AIRAS (Tanaka et al., 2025).

REFERENCES

Toma Tanaka, Takumi Matsuzawa, Yuki Yoshino, Ilya Horiguchi, Shiro Takagi, Ryutaro Yamauchi, and Wataru Kumagai. AIRAS, 2025. URL https://github.com/airas-org/airas.